

Erlang and Distributed Computing

Colin Morris · Qiyu Zhu

December 2010



Questions

Is Erlang an appropriate tool for high-performance computing?

Does Erlang really make parallel programming and fault-tolerance easy?

Can we characterize problems that are easy and difficult to make fault-tolerant?



What is Erlang?

A concurrent, functional programming language.

The result of Ericsson researchers seeking to design a programming language “that led to the shortest and most beautiful programs closest to the level of formal specifications.”



Erlang by Example

C

```
X=1;  
X=2;
```

Math

```
X = 1  
X = 2 X
```

Like math, Erlang is single assignment. Once bound, a variable cannot be assigned to again. This is one way Erlang reduces side effects.

Erlang

```
X=1.  
X=2.  
** exception error:  
no match of right  
hand side value 2
```



Erlang by Example pt. 2

C

```
int factorial(int n) {  
    int result = 1;  
    for (int i=1; i<=n; i++){  
        result *= i;  
    }  
    return result;  
}
```

Math

$$n! = \begin{cases} 1 & \text{if } n = 0, \\ (n-1)! \times n & \text{if } n > 0. \end{cases}$$

Erlang

```
fact(0) -> 1;  
fact(n) -> n*fact(n-1).
```

Erlang
performs
repetition only
by recursion.
No loops.

A Refinement

```
fact(n) -> fact(n, 1) .
```

```
fact(0, Acc) -> Acc;
```

```
fact(n, Acc) -> fact(n-1, n*Acc) .
```

We use **tail recursion** to avoid growing the stack. We keep a running sum in an accumulator variable *Acc*, and never have to pass control back up to the calling function.



Erlang by Example pt. 3

C

```
int j = 0;
for (int i=0; i<sizeof(S)/sizeof(*S); i++) {
    if (S[i]%2==0) {
        E[j] = S[i];
        j++;
    }
}
```

Math

$$E = \{x \in S : 2 \mid x\}$$

Erlang

```
E = [X || X <- S, X rem 2 == 0]
```



Concurrency in Erlang

Erlang uses asynchronous message passing as its only form of interprocess communication.

No shared memory.

Actor model.



Fault Tolerance

From the beginning, Erlang was designed to be fault tolerant. In the face of hardware or software errors, a fault-tolerant system should march on.

Erlang encourages “programming for the correct case”.

If something goes wrong, the process can just fail and be restarted.

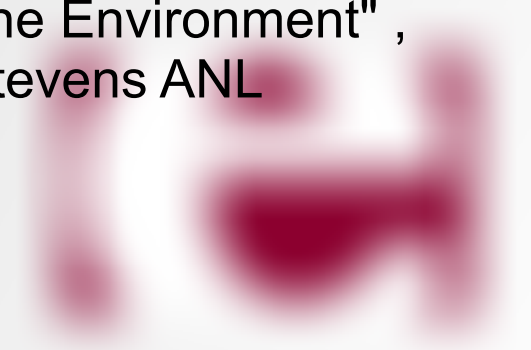
Or can it?



Why does fault tolerance matter?

Modern PCs may run for weeks without needing rebooting. Today's supercomputers often run for only a few days before rebooting, because of their complexity and their thousands of processors. Exascale systems will be even more complex and have millions of processors.... application software will not be able to rely on checkpoint/restart to cope with faults since a new fault is likely to occur before the application can be restarted. For exascale systems, new fault tolerance paradigms will need to be developed and integrated into both existing and new applications.

"Modeling and Simulation at the Exascale for Energy and the Environment" ,
2007: Horst Simon LBNL, Thomas Zacharia ONRL, Rick Stevens ANL



What *doesn't* Erlang do?

Erlang has limited support for arrays, and no support for multi-dimensional arrays.

Erlang is not fast. We found it to be about 50–100 times slower than C for numerical computation.

No numerical libraries.



Case studies



Cowichan Problems

A set of simple computation problems with same properties as real HPC applications

Tests the programmer's ability and the programming environment's capability

Illustrates the practicality of new programming paradigms and parallelization

<http://code.google.com/p/cowichan/>




Case study: Mandelbrot set

Problem: Generate the Mandelbrot Set for a specified region of the complex plane.

Example:

`mandel(10, 10, -1.5, -1.5, 3, 3).`



```
{2, 2, 2, 2, 2, 2, 2, 2, 2, 2},
{3, 3, 3, 4, 12, 4, 2, 2, 2, 2},
{3, 3, 4, 6, 150, 6, 4, 2, 2, 2},
{4, 5, 7, 150, 150, 150, 6, 3, 2, 2},
{9, 150, 22, 150, 150, 150, 6, 3, 2, 2},
{9, 150, 22, 150, 150, 150, 6, 3, 2, 2},
{4, 5, 7, 150, 150, 150, 6, 3, 2, 2},
{3, 3, 4, 6, 150, 6, 4, 2, 2, 2},
{3, 3, 3, 4, 12, 4, 2, 2, 2, 2},
{2, 2, 2, 2, 2, 2, 2, 2, 2, 2},
```

Mandelbrot set

Easy to implement

Embarrassingly parallel

Only fork-join communication needed

Easy to recover from failures

Uneven workload



The difficulty

Erlang has no multi-dimensional arrays

We wrote a 2D array library using lists of lists or arrays of arrays

Erlang

```
lists:sublist(A, j-1) ++  
[lists:sublist(lists:nth(j, A),  
k-1) ++ [n] ++  
lists:sublist(lists:nth(j, A),  
k+1, length(A))]  
++ lists:sublist(A, j+1,  
length(A)).
```

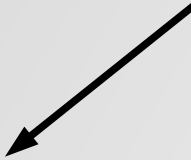
C

$A[j][k] = n$



Erlang

This is supposed to be one line




```
lists:sublist(A, j-1) ++  
[lists:sublist(lists:nth(j, A),  
k-1) ++ [n] ++  
lists:sublist(lists:nth(j, A),  
k+1, length(A))]  
++ lists:sublist(A, j+1,  
length(A)).
```

The easy

Embarrassingly parallel: we can just apply a parallel map.

Easy to make fault tolerant: little sense of state. Any one worker process could die and be restarted without affecting any other worker processes.

In this sense, of all the problems we looked at, Mandelbrot fit closest with the functional paradigm.



Fault tolerance in Mandelbrot

```
parallel_mandel(Nrows, Ncols, X0, Y0, Dx, Dy) ->
  Mat = init_matrix(Nrows, Ncols, X0, Y0, Dx, Dy),
  pmap(fun mandel:mandel_number/1, Mat).

% Simple parallel map function
pmap(Fun, Arr) -> Parent = self(),
  Pids = lists:map(fun(Ele) ->
    spawn(fun() -> pmap_f(Parent, Fun, Ele) end) end, Arr),
  pmap_gather(lists:zip(Pids, Arr), Fun).

pmap_f(Parent, Fun, Element) ->
  Parent ! {self(), lists:map(Fun, Element)}.

pmap_gather([{Pid, Elem}|T], Fun) ->
  receive
    {Pid, Ret} -> [Ret|pmap_gather(T)]
  after 2000 ->
    spawn(fun() -> pmap_f(self(), Fun, Elem) end),
    pmap_gather(T ++ [{Pid, Elem}], Fun)
  end;
pmap_gather([], _) -> [].
```

Case study: Game of Life

Small amount of computation per cell

Uniform workload

Peer-to-peer communication

But

Reading neighbor cells is awkward

Hard to implement fault tolerance efficiently





Strategy

Divide grid into vertical slices of equal width, with one process per slice

Between each iteration, processes send and receive their extreme columns with neighbors



Example



Case study: Sample sort

Similar to Game of Life, but even more communication intensive


Algorithm is divided into distinct stages, allowing us to at least checkpoint between stages, if not during



Performance



Compared to C




Strong and weak scaling

Strong scaling: Fixed problem size, variable number of processors. Expect reciprocally decreasing time.


Weak scaling: Amount of computation is made proportional to number of processors. Expect constant time.




Mandelbrot, strong scaling, 1 node




Mandelbrot, strong scaling




Mandelbrot, strong scaling, speedup




Mandelbrot, weak scaling




Life, strong scaling, 10 nodes



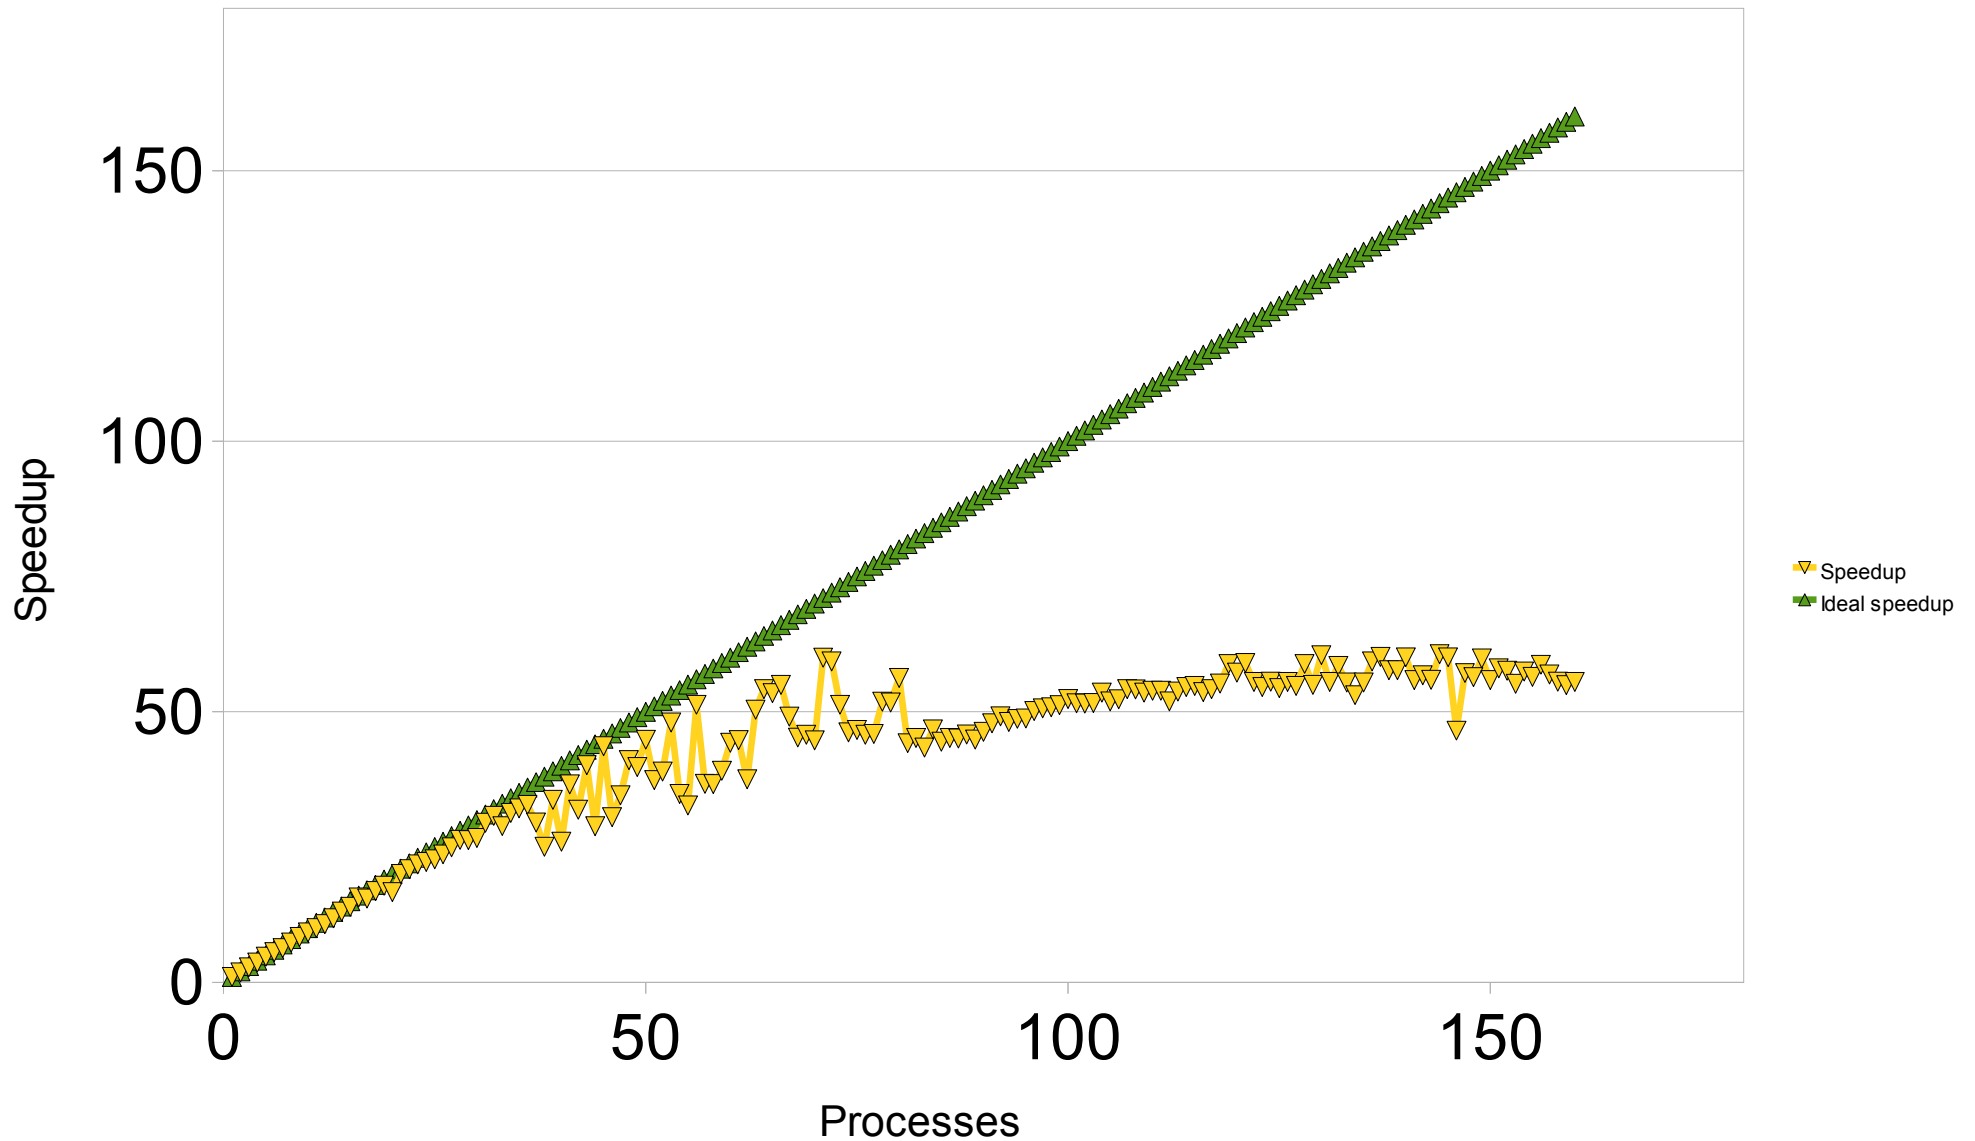
Life, strong scaling, speedup



Life, weak scaling



Simplified Mandelbrot, strong scaling



Erlang and C

Can use Erlang and C to do parallel computation

Erlang handles process control signals

C code does the heavy computation



Acknowledgements

Erlang package (open source)

Mentors: Greg Wilson, Jonathan Dursi

SciNet General Purpose Cluster (GPC)



Summary

Erlang is not a get-out-of-jail-free card

Parallelization and fault tolerance require planning, and there is no general strategy



Q&A

